



Introducing OpenMP

Jonathan Ferland


Scientific computing analyst

OpenMP objectives

- Standardization : available on all shared memory architectures.
- Simple interface : reduced instructions set (3 to 4) usually sufficient for parallelization.
- Ease of use : progressive integration possible.
- Portable: supported in many languages :
Fortran (77,90,95), C (C90,C99) and C++



OpenMP model



OpenMP model cont.

- Master
 - Executes sequential part of the program
 - Participates in parallel region
 - Exists for program lifetime
- Workers/Children
 - In general, all of them execute the same task on a different part of the data
- Number of threads is independent on the number of processors.



First code

- To use OpenMP, the following is needed:

#include <omp.h> (C)

use omp_lib (Fortran 90,95)

include 'omp_lib.h' (Fortran 77)

- A parallel region is defined by a delimiter (known by the compiler) and a “parallel” directive.

#pragma omp parallel {} (C)

!\$omp parallel (Fortran)

!\$omp end parallel



First code cont.

- It might be useful to know a thread rank or the total number of threads inside the parallel region.

total = omp_get_num_threads()

rank = omp_get_thread_num()

- Variable scope is important.

private(rank) shared(total)

- Let's now look at a simple “hello world” example using OpenMP.



Hello world! (C)

```
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif

int main (int argc, char ** argv)
{
    int total =1, rank=0;

    #pragma omp parallel private(rank) shared(total)
    {
#ifdef _OPENMP
        total = omp_get_num_threads();
        rank = omp_get_thread_num();
#endif
        printf("Hello from thread %d in a group of %d threads\n",rank,total);
    }
}
```



Hello world! (Fortran)

```
program hello
!$ use omp_lib

integer total, rank

total=1
rank=0
!$omp parallel private(rank) shared(total)
!$  total = omp_get_num_threads()
!$  rank = omp_get_thread_num()
print *, "Hello from thread ", rank, " in a group of ", total, " threads"
!$omp end parallel
end program hello
```



Observations

- Must include `omp.h` or use “use `omp_lib`”.
- Thread rank ranges from 0 up to total-1.
- The macro `_OPENMP` is defined when compiling with `-openmp` in C to exclude/include code. Sentinel syntax in Fortran allows for conditional inclusion of code.
 - Sentinel can be : `!$` or `C$` or `*$` (fixed format)

`!$` (free form format)
- All threads inside a parallel region are doing the same work, but using different data.



Compilation and execution

- To compile an openMP application, extra options are needed.

```
ifort/icc hello.[f90,c] -o prog_hello -openmp
```

- Total number of threads inside parallel region can be set via environment variable.

```
setenv OMP_NUM_THREADS 4
```

```
./prog_hello
```

```
Hello from thread 0 in a group of 4 threads
```

```
Hello from thread 1 in a group of 4 threads
```

```
Hello from thread 3 in a group of 4 threads
```

```
Hello from thread 2 in a group of 4 threads
```



Work sharing

- To assign given task to a specific thread.
- To use do/for loops, one must use directives:

#pragma omp for (C)

!\$omp do

!\$omp end do (Fortran)

- By default, the loop's iterations are equally distributed on all threads.



thread 1

thread 2

thread 3

thread 4



Loop (C)

```
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif
#define N 20

int main (int argc, char ** argv)
{
    int i, rank=0;

    #pragma omp parallel private(rank)
    {
        #ifdef _OPENMP
            rank = omp_get_thread_num();
        #endif
        #pragma omp for
        for (i=0; i<N; i++)
            printf("Element %d done by thread %d \n", i, rank);
    }
}
```



Loop (Fortran)

```
program loop
!$ use omp_lib

integer i, rank, N
N=20
rank=0
!$omp parallel private(rank)
!$  rank = omp_get_thread_num()
!$omp do
do i=1, N
  print *, "Element",i," done by thread ",rank
end do
!$omp end parallel
end program loop
```



Observations

- Loop index is private.
- the do/for loop must immediately follow the do/for sentinel declaration.
- Must be inside a parallel region.
- There is a short form of the do/for sentinel.

#pragma omp parallel for (C)

!\$omp parallel do

!\$omp end parallel do (Fortran)

- The sentinel directive “**!\$omp end do**” is optional.



Work sharing cont.

- Each thread can execute its own work, which can be different from what other threads are doing.
- The following is used:

C

```
#pragma omp sections  
{  
#pragma omp section  
{ ... }  
#pragma omp section  
{ ... }  
}
```

Fortran

```
!$omp sections  
!$omp section  
...  
!$omp section  
...  
!$omp end sections
```



Sections (C)

```
...
int main (int argc, char ** argv)
{
    int rank=0;

#pragma omp parallel private(rank)
    {
#ifdef _OPENMP
        rank = omp_get_thread_num();
#endif
#pragma omp sections
    {
#pragma omp section
        printf("Section 1 is executed by thread %d\n", rank);
#pragma omp section
        printf("Section 2 is executed by thread %d\n", rank);
    }
}
```



Sections (Fortran)

...

integer rank

rank=0

!\$omp parallel private(rank)

!\$ rank = omp_get_thread_num()

!\$omp sections

!\$omp section

print *, "Section 1 is executed by thread" , rank

!\$omp section

print *, "Section 2 is executed by thread" , rank

!\$omp end sections

!\$omp end parallel

...



Observations

- Each section is executed by a different thread.
- Load balancing between sections is important.
- There is a short form for the sections sentinel:

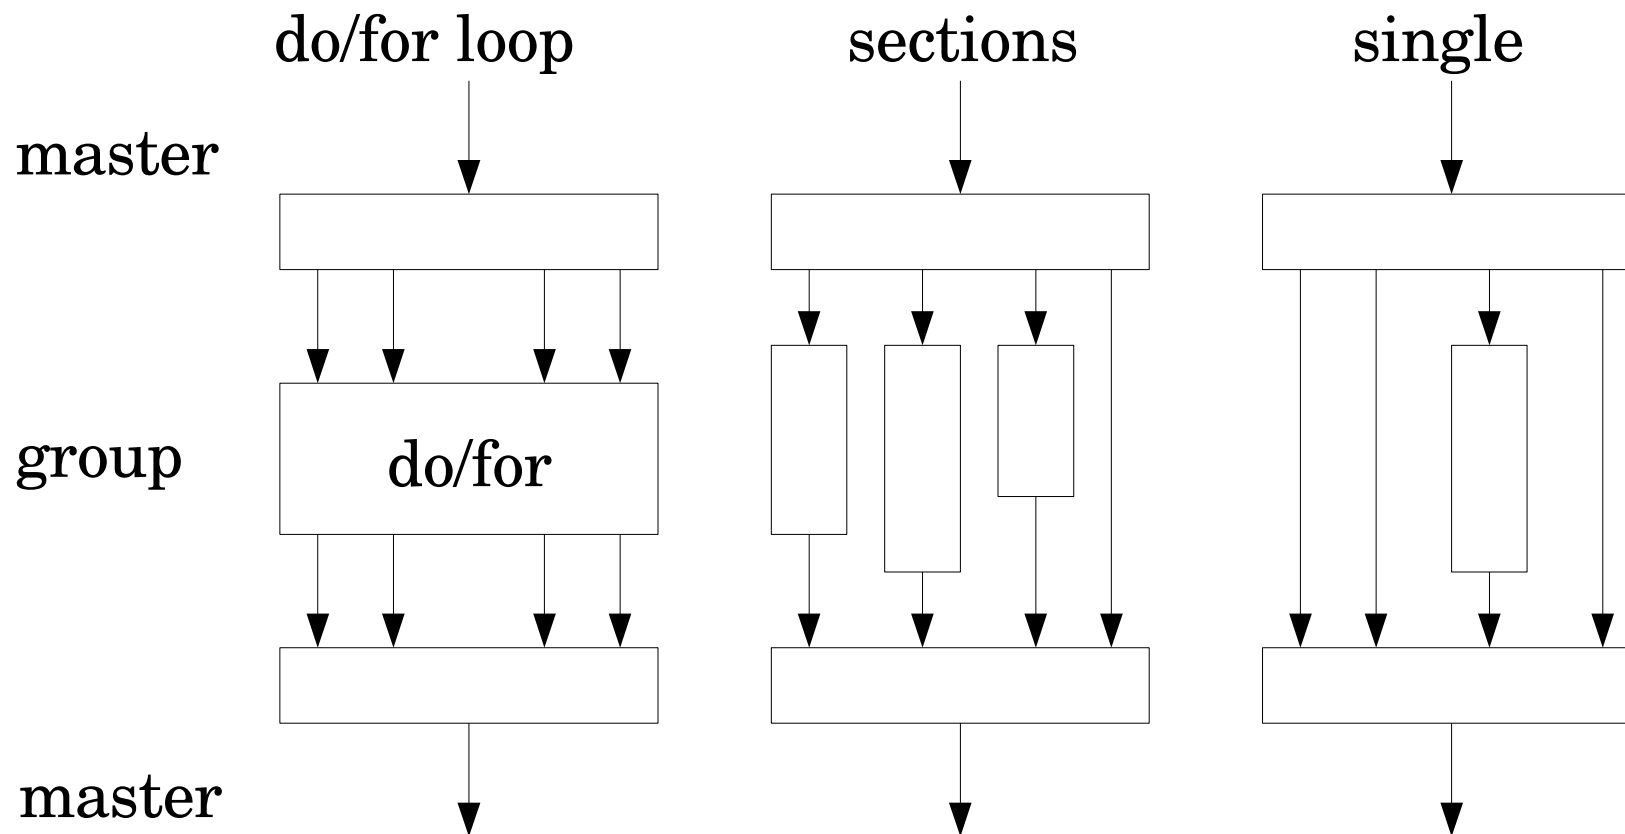
#pragma omp parallel sections (C)

!\$omp parallel sections

!\$omp end parallel sections (Fortran)



Work sharing cont.



Reduction

- To carry out an operation (sum, multiplication, etc.) involving all threads. Some logical operation are also supported.
- In addition to shared and private, a variable's scope can be “reduction”.



Naive reduction (C)

```
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif
#define N 20

int main (int argc, char ** argv)
{
    int i, sum = 0, val =0, elements[N];
    #pragma omp parallel for shared(elements)
    for (i=0; i<N; i++)  elements[i] = i+1;

    #pragma omp parallel private(val) shared (elements,sum)
    #pragma omp for
    for (i=0; i<N; i++)
        val += elements[i];
    #pragma omp critical
    sum += val;

    printf ("The sum is %d\n", sum);
}
```



Naive reduction (Fortran)

```
program reduction
!$ use omp_lib
integer i, sum, val, elements(20)
sum = 0
val = 0
!$omp parallel do shared(elements)
do i=1,20
  elements(i) = i
end do
!$omp parallel private(val) shared(elements,sum)
!$omp do
do i=1,20
  val = val + elements(i)
end do
!$omp end do
!$omp critical
sum = sum+val
!$omp end critical
!$omp end parallel
print *, "The sum is ", sum
end program reduction
```



Reduction (C)

```
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif
#define N 20

int main (int argc, char ** argv)
{
    int i, sum = 0, elements[N];
    #pragma omp parallel for shared(elements)
    for (i=0; i<N; i++)
        elements[i] = i+1;

    #pragma omp parallel for shared (elements) reduction (+:sum)
    for (i=0; i<N; i++)
        sum += elements[i];

    printf ("The sum is %d\n", sum);
}
```



Reduction (Fortran)

```
program reduction
!$ use omp_lib

integer i, sum, elements(20)
sum = 0
!$omp parallel do shared(elements)
do i=1,20
  elements(i) = i
end do

!$omp parallel do shared(elements) reduction(+:sum)
do i=1,20
  sum = sum + elements(i)
end do
!$omp end parallel do

print *, "The sum is ", sum
end program reduction
```




Observations

- Available operators are +, -, * and some boolean operators. In Fortran min and max operator are also available for reduction.
- Multiple reduction operations can be realized in the same loop.
- Results may differ from those obtained during serial execution due to round-off errors.
- Reductions also work with section constructions.
- Initializing tables in parallel is better (on Altix system, the performance is better).



Altix system



Multiple hosts are interconnected by a fast network and are configured in a way that makes them look like a single big node. In such a system, memory access isn't uniform. “dplace” is a tool to pin processes on a given underlying node in order to have fast memory accesses.

Performance

```
program reduction
use iflport
!$ use omp_lib
integer i, sum, elements(1000000)
real*8 begin, end

do i=1,1000000
  elements(i) = irand()
end do

begin = dclock()
!$omp parallel do shared(elements) reduction(+:sum)
do i=1,1000000
  sum = sum + elements(i)
end do
!$omp end parallel do

end = dclock()
print *, "The sum is ", sum, " and took",(end-begin)*1000.0d0," m sec"
end program reduction
```



Performance cont.

- OpenMP is not always a good solution. There is an overhead when using it.

Size	Serial	OpenMP(4)
1 000 000	0.5 ms	3.4 ms
100 000 000	115 ms	55 ms

- OpenMP might be a good choice if the amount of work to do in each thread is large or if the total time spent inside one thread is high.
- It is very important to profile your application in order to decide where OpenMP might be useful.




References

- <http://www.openmp.org>
- *Parallel Programming in Openmp*, Chandra, et al., Morgan Kaufmann publisher, 2001.
- <https://computing.llnl.gov/tutorials/openMP>



Hybrid MPI/OpenMP

- The inter-node communication uses MPI from the openMP master thread on each node.
- Calculations are executed on multiple openMP threads on every node.



Pthreads

- Posix threads allow you to spawn a new concurrent process flow. It defines ways to schedule, synchronize, etc.
- A standard available on all platforms (C/C++).
- Very flexible.
- Hard to debug and to design on large system.
- Pthreads: Programming with POSIX Threads, David R. Butenhof, Addison-Wesley Professional computing series, 1997.



Others

Compilers

- UPC (*Unified parallel C*): <http://www.gwu.edu/~upc/>
- Co-array Fortran: <http://www.co-array.org/>

Libraries

- SHMEM: `man intro_shmem` (Altix)
- Global arrays: <http://www.emsl.pnl.gov/docs/global/>

